

CSC 722 - Machine Learning Fundamentals

CNN Team Project

Team Members:

1. Dheeraj Avadhutha - 101129707
2. Pranav Reddy Dareddy - 101134583
3. Vamsi Yadala - 101130145
4. Sri Mukund Kadiyala - 101134587
5. Chenchaiiah Mekalathuru - 101165473
6. Sai Krishna Yadav Madiboyena - 101132195
7. Shanmukh Sai Madhu - 101162913
8. Pavan Kumar Balli - 101136776
9. Sai Rajesh Chittavarjula - 101179899
10. Rohit Mareddy - 101134589

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
```

```
# Create a Convolutional Neural Network (CNN) model
cnn_model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    MaxPooling2D((2, 2)),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Conv2D(64, (3, 3), activation='relu'),
    Flatten(),
    Dense(64, activation='relu'),
    Dense(10, activation='softmax')
])
```

```
import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.utils import to_categorical
```

```
# Load and preprocess the MNIST dataset
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
train_images = train_images.reshape(-1, 28, 28, 1) / 255.0
test_images = test_images.reshape(-1, 28, 28, 1) / 255.0
train_labels = to_categorical(train_labels, num_classes=10)
test_labels = to_categorical(test_labels, num_classes=10)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
11490434/11490434 [=====] - 0s 0us/step

```
# Display model summary
cnn_model.summary()
```

```
from tensorflow.keras.layers import MaxPooling2D
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_2 (Conv2D)	(None, 3, 3, 64)	36928
flatten (Flatten)	(None, 576)	0
dense (Dense)	(None, 64)	36928
dense_1 (Dense)	(None, 10)	650

```

=====
Total params: 93322 (364.54 KB)
Trainable params: 93322 (364.54 KB)
Non-trainable params: 0 (0.00 Byte)
=====

```

```

# Example of max pooling layer
max_pool_layer = MaxPooling2D(pool_size=(2, 2))

```

```

from tensorflow.keras.layers import Dense

```

```

# Fully connected layer followed by softmax
fc_layer = Dense(64, activation='relu')
softmax_layer = Dense(10, activation='softmax')

```

```

# Compile and train the CNN model
cnn_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = cnn_model.fit(train_images, train_labels, epochs=10, batch_size=128, validation_data=(test_images, test_labels))

```

```

Epoch 1/10
469/469 [=====] - 57s 118ms/step - loss: 0.2497 - accuracy: 0.9241 - val_loss: 0.0660 - val_accuracy: 0.9804
Epoch 2/10
469/469 [=====] - 48s 102ms/step - loss: 0.0627 - accuracy: 0.9804 - val_loss: 0.0514 - val_accuracy: 0.9865
Epoch 3/10
469/469 [=====] - 51s 108ms/step - loss: 0.0426 - accuracy: 0.9865 - val_loss: 0.0347 - val_accuracy: 0.9898
Epoch 4/10
469/469 [=====] - 49s 104ms/step - loss: 0.0334 - accuracy: 0.9898 - val_loss: 0.0303 - val_accuracy: 0.9910
Epoch 5/10
469/469 [=====] - 49s 104ms/step - loss: 0.0278 - accuracy: 0.9910 - val_loss: 0.0373 - val_accuracy: 0.9931
Epoch 6/10
469/469 [=====] - 49s 104ms/step - loss: 0.0223 - accuracy: 0.9931 - val_loss: 0.0336 - val_accuracy: 0.9937
Epoch 7/10
469/469 [=====] - 50s 107ms/step - loss: 0.0193 - accuracy: 0.9937 - val_loss: 0.0293 - val_accuracy: 0.9952
Epoch 8/10
469/469 [=====] - 53s 113ms/step - loss: 0.0150 - accuracy: 0.9952 - val_loss: 0.0289 - val_accuracy: 0.9956
Epoch 9/10
469/469 [=====] - 50s 107ms/step - loss: 0.0136 - accuracy: 0.9956 - val_loss: 0.0319 - val_accuracy: 0.9969
Epoch 10/10
469/469 [=====] - 47s 101ms/step - loss: 0.0098 - accuracy: 0.9969 - val_loss: 0.0284 - val_accuracy: 0.9969

```

```

# Evaluate the model
test_loss, test_accuracy = cnn_model.evaluate(test_images, test_labels)
print(f'Test Accuracy: {test_accuracy}')

```

```

from sklearn.model_selection import KFold
from sklearn.metrics import confusion_matrix

```

```

313/313 [=====] - 3s 9ms/step - loss: 0.0284 - accuracy: 0.9919
Test Accuracy: 0.9919000267982483

```

```

# Use KFold for cross-validation
kf = KFold(n_splits=5, shuffle=True)
for train_index, test_index in kf.split(train_images):
    X_train_fold, X_val_fold = train_images[train_index], train_images[test_index]
    y_train_fold, y_val_fold = train_labels[train_index], train_labels[test_index]

```

```

# Train and evaluate model for each fold
cnn_model.fit(X_train_fold, y_train_fold)
val_loss, val_accuracy = cnn_model.evaluate(X_val_fold, y_val_fold)
print(f'Validation Accuracy: {val_accuracy}')

```

```

# Generate confusion matrix
predictions = cnn_model.predict(X_val_fold)
cm = confusion_matrix(y_val_fold.argmax(axis=1), predictions.argmax(axis=1))
print(cm)

```

```

1500/1500 [=====] - 45s 30ms/step - loss: 0.0325 - accuracy: 0.9890
375/375 [=====] - 3s 9ms/step - loss: 0.0274 - accuracy: 0.9908
Validation Accuracy: 0.9908333420753479
375/375 [=====] - 4s 10ms/step
[[1155  0  0  0  1  0  1  0  2  1]
 [  0 1401  0  0  0  0  0  7  4  0]
 [  2  1 1151  1  0  0  0  5  1  0]
 [  2  0  2 1229  0  6  0  2  2  2]
 [  0  1  0  0 1094  0  1  4  6 13]]

```

```

[ 0 0 1 1 0 1101 1 0 1 2]
[ 3 2 1 0 1 0 1159 0 2 0]
[ 0 1 3 0 0 0 0 1220 2 0]
[ 2 0 0 0 0 4 1 1 1169 2]
[ 0 0 0 1 2 2 0 3 4 1211]]
1500/1500 [=====] - 44s 29ms/step - loss: 0.0209 - accuracy: 0.9934
375/375 [=====] - 5s 12ms/step - loss: 0.0180 - accuracy: 0.9939
Validation Accuracy: 0.9939166903495789
375/375 [=====] - 3s 9ms/step
[[1191 0 0 0 0 0 0 0 0 1]
[ 0 1340 1 0 1 0 1 5 0 1]
[ 0 1 1155 2 0 0 0 3 0 1]
[ 1 0 4 1238 0 4 0 1 2 1]
[ 0 1 0 0 1155 0 0 2 0 10]
[ 0 0 0 1 0 1058 3 1 1 1]
[ 3 0 0 0 0 1 1191 0 0 0]
[ 0 0 1 0 0 0 0 1273 0 1]
[ 1 1 0 0 3 2 0 0 1175 8]
[ 0 0 0 0 2 0 0 0 0 1151]]
1500/1500 [=====] - 44s 30ms/step - loss: 0.0177 - accuracy: 0.9941
375/375 [=====] - 3s 9ms/step - loss: 0.0206 - accuracy: 0.9937
Validation Accuracy: 0.9937499761581421
375/375 [=====] - 3s 9ms/step
[[1187 0 2 1 0 1 0 0 0 0]
[ 0 1256 0 0 0 0 0 3 0 0]
[ 0 2 1189 10 0 0 0 3 0 0]
[ 0 1 0 1222 0 1 0 1 0 0]
[ 0 1 0 0 1194 0 2 4 0 1]
[ 0 0 0 3 0 1121 0 0 0 0]
[ 0 1 0 1 1 5 1183 0 8 0]
[ 0 0 0 0 0 0 0 1258 0 0]
[ 0 0 2 2 0 0 0 0 1128 1]
[ 0 2 0 1 3 1 0 11 0 1187]]
1500/1500 [=====] - 42s 28ms/step - loss: 0.0157 - accuracy: 0.9951
375/375 [=====] - 4s 10ms/step - loss: 0.0079 - accuracy: 0.9974
Validation Accuracy: 0.9974166750907898
375/375 [=====] - 3s 9ms/step
[[1174 0 0 0 0 0 2 0 1 0]
[ 0 1394 0 0 1 0 0 0 1 0]
[ 0 0 1213 0 1 0 0 0 0 0]
[ 0 0 1 1189 0 0 0 0 0 0]
[ 0 0 0 0 1217 0 1 0 0 1]
[ 1 0 0 1 0 1062 1 0 1 0]
[ 1 0 0 0 1 1 1137 0 0 0]
[ 0 5 1 1 1 0 0 1274 0 1]
[ 0 0 0 0 0 0 1 0 1115 1]
[ 0 2 0 0 1 0 0 1 0 1194]]
1500/1500 [=====] - 44s 29ms/step - loss: 0.0111 - accuracy: 0.9963
375/375 [=====] - 3s 9ms/step - loss: 0.0064 - accuracy: 0.9976

```